

Building Generally Capable AI Agents with MineDojo

<https://developer.nvidia.com/blog/building-generally-capable-ai-agents-with-minedojo/>

Admin

- **Draft feedback from TAs going out today**
- Marks later this week ...
- **Don't forget weekly project standups!**

The Data of RL


Imagine you developed an new algorithm

- One of the primary ways to understand and evaluate your new idea is via experiments
- There are many things you might want to know:
 - Is my implementation correct?
 - Does the method converge to the correct thing?
 - How does the performance vary as a function of initialization, hyper parameters, and design choices?
 - What are the limitations of the idea?
 - Lastly, if it is better in some measurable, reliable, relevant way?

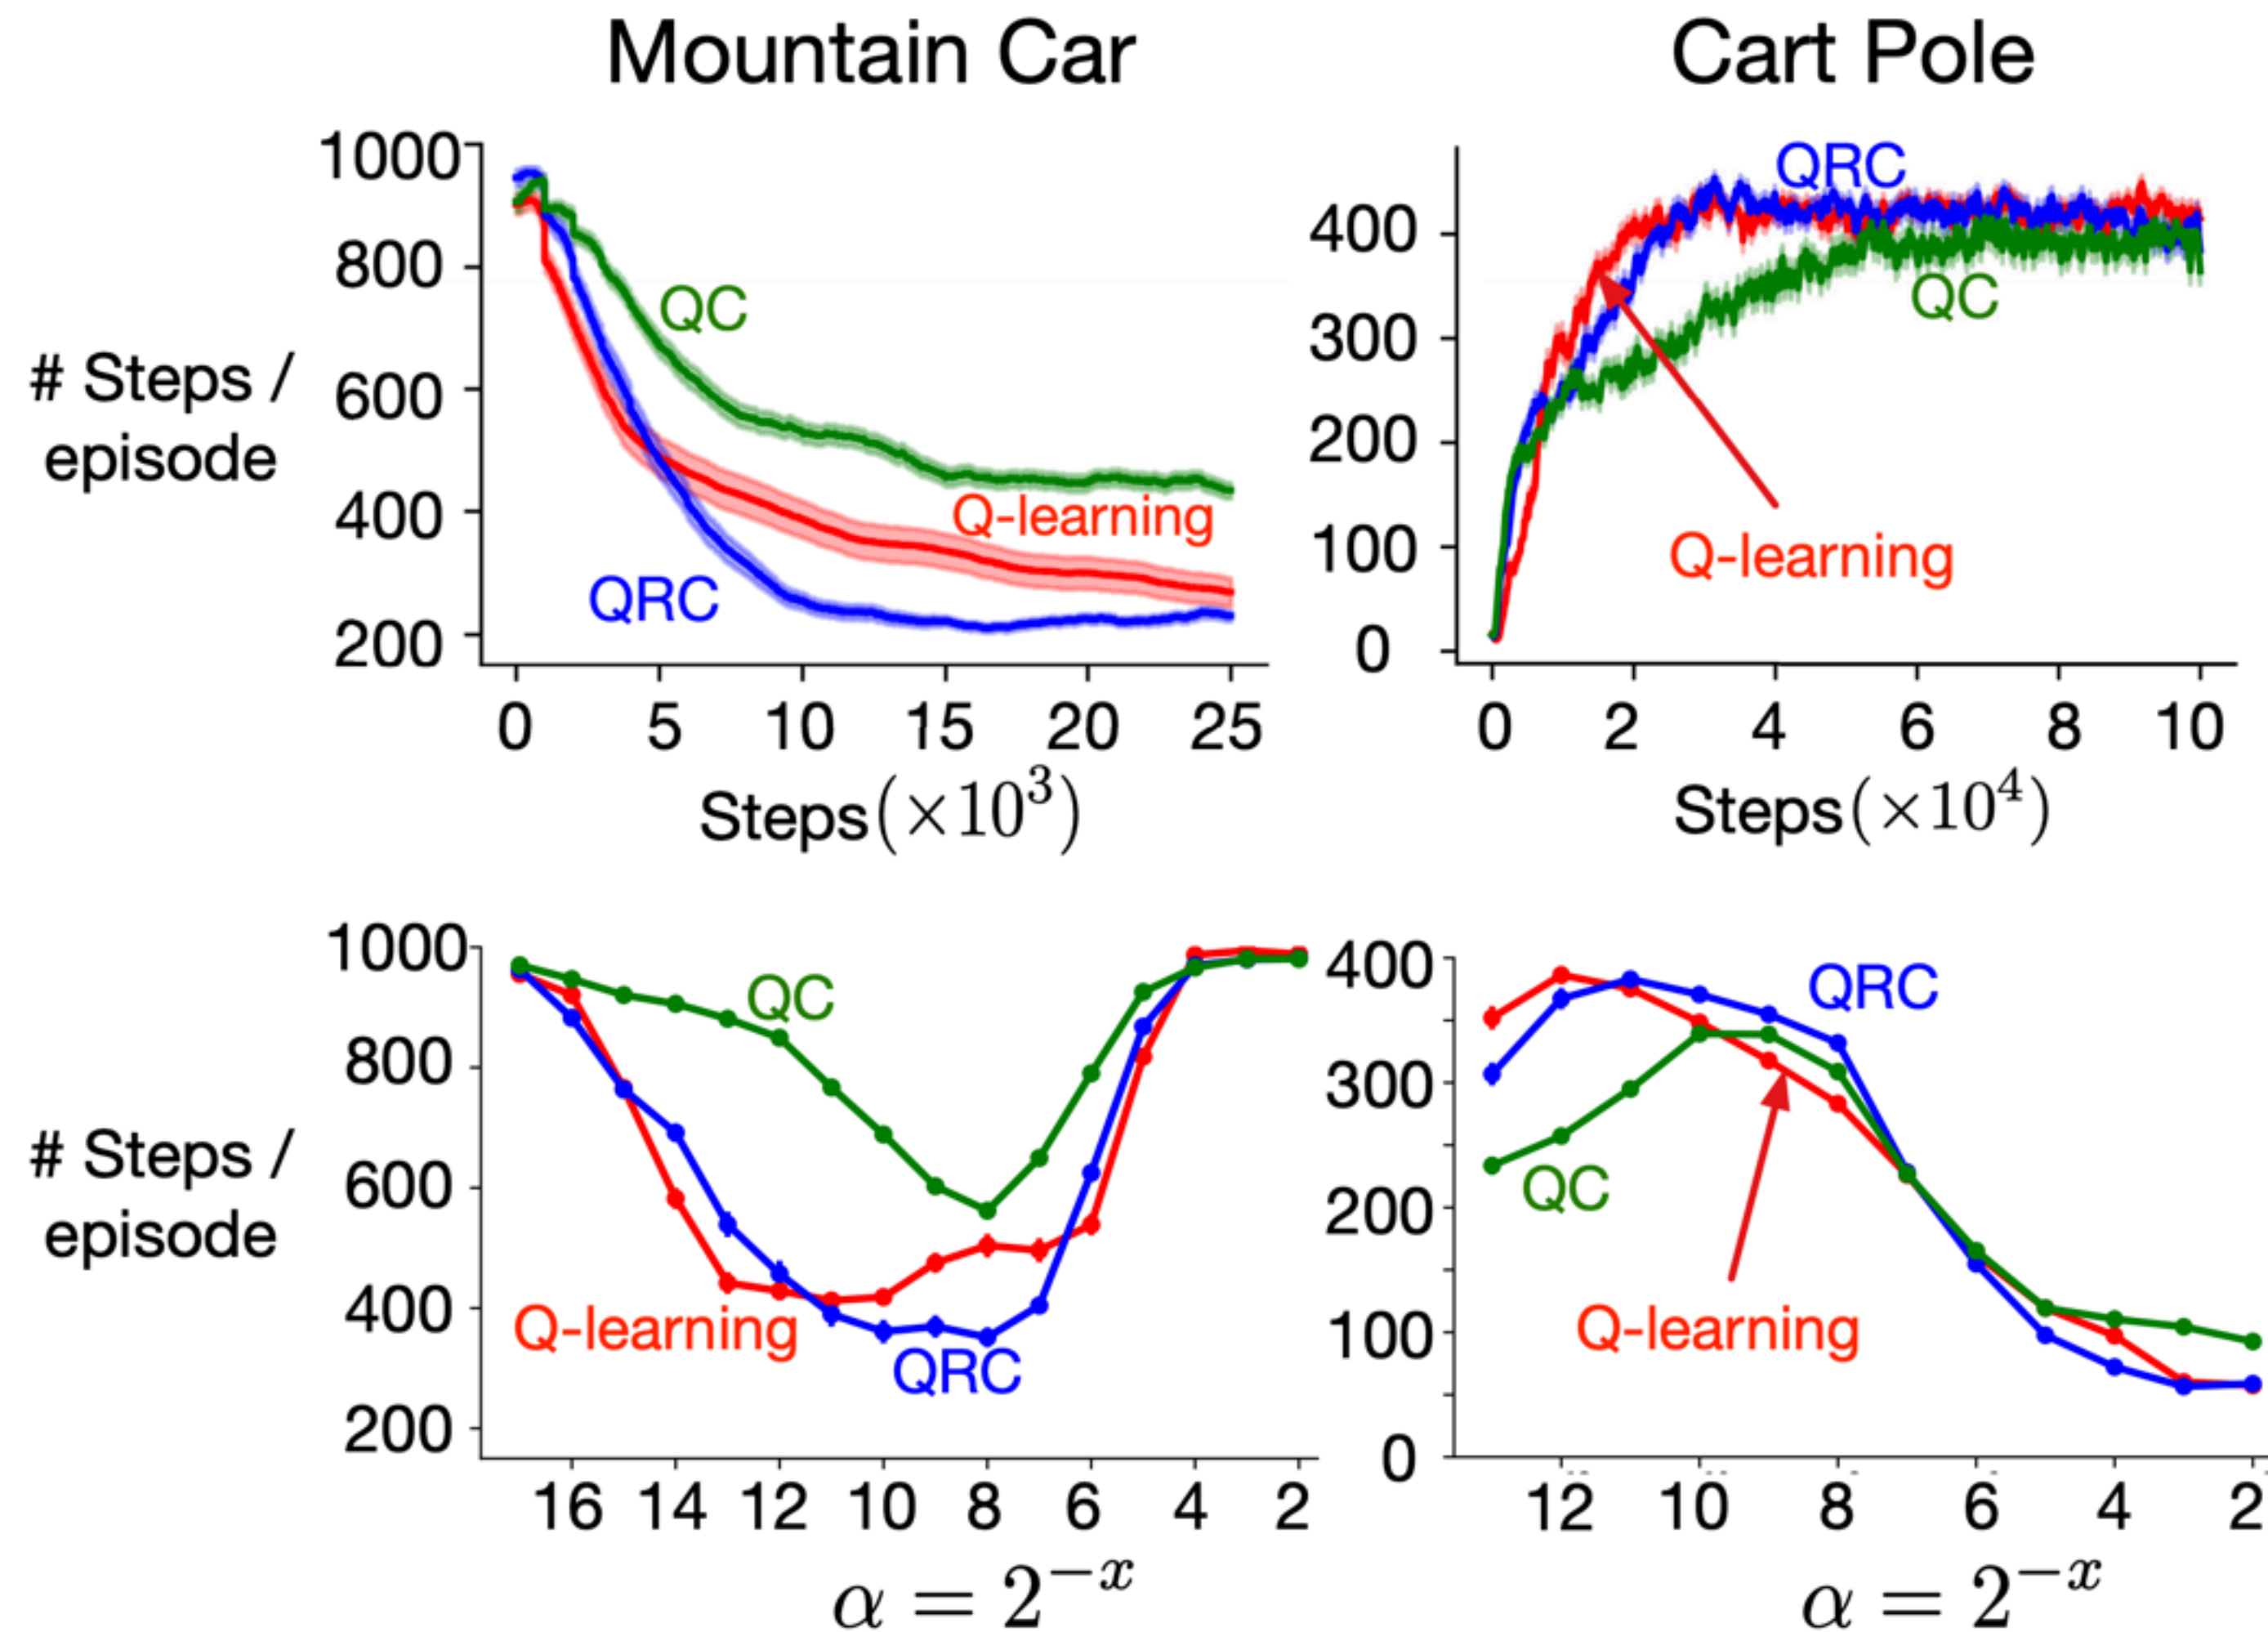
Start with the problem

- Common failure:
 - Spend time developing a new approach, and adjusting your experiments to illustrate the new approach works and works well
 - Someone points out a missing baseline or alternative approach
 - The baseline is better than all the other algorithms tested
- Alternative strategy:
 - Start with the open problem
 - Show that baselines fail or have some important limitations

Example: step-size adaption



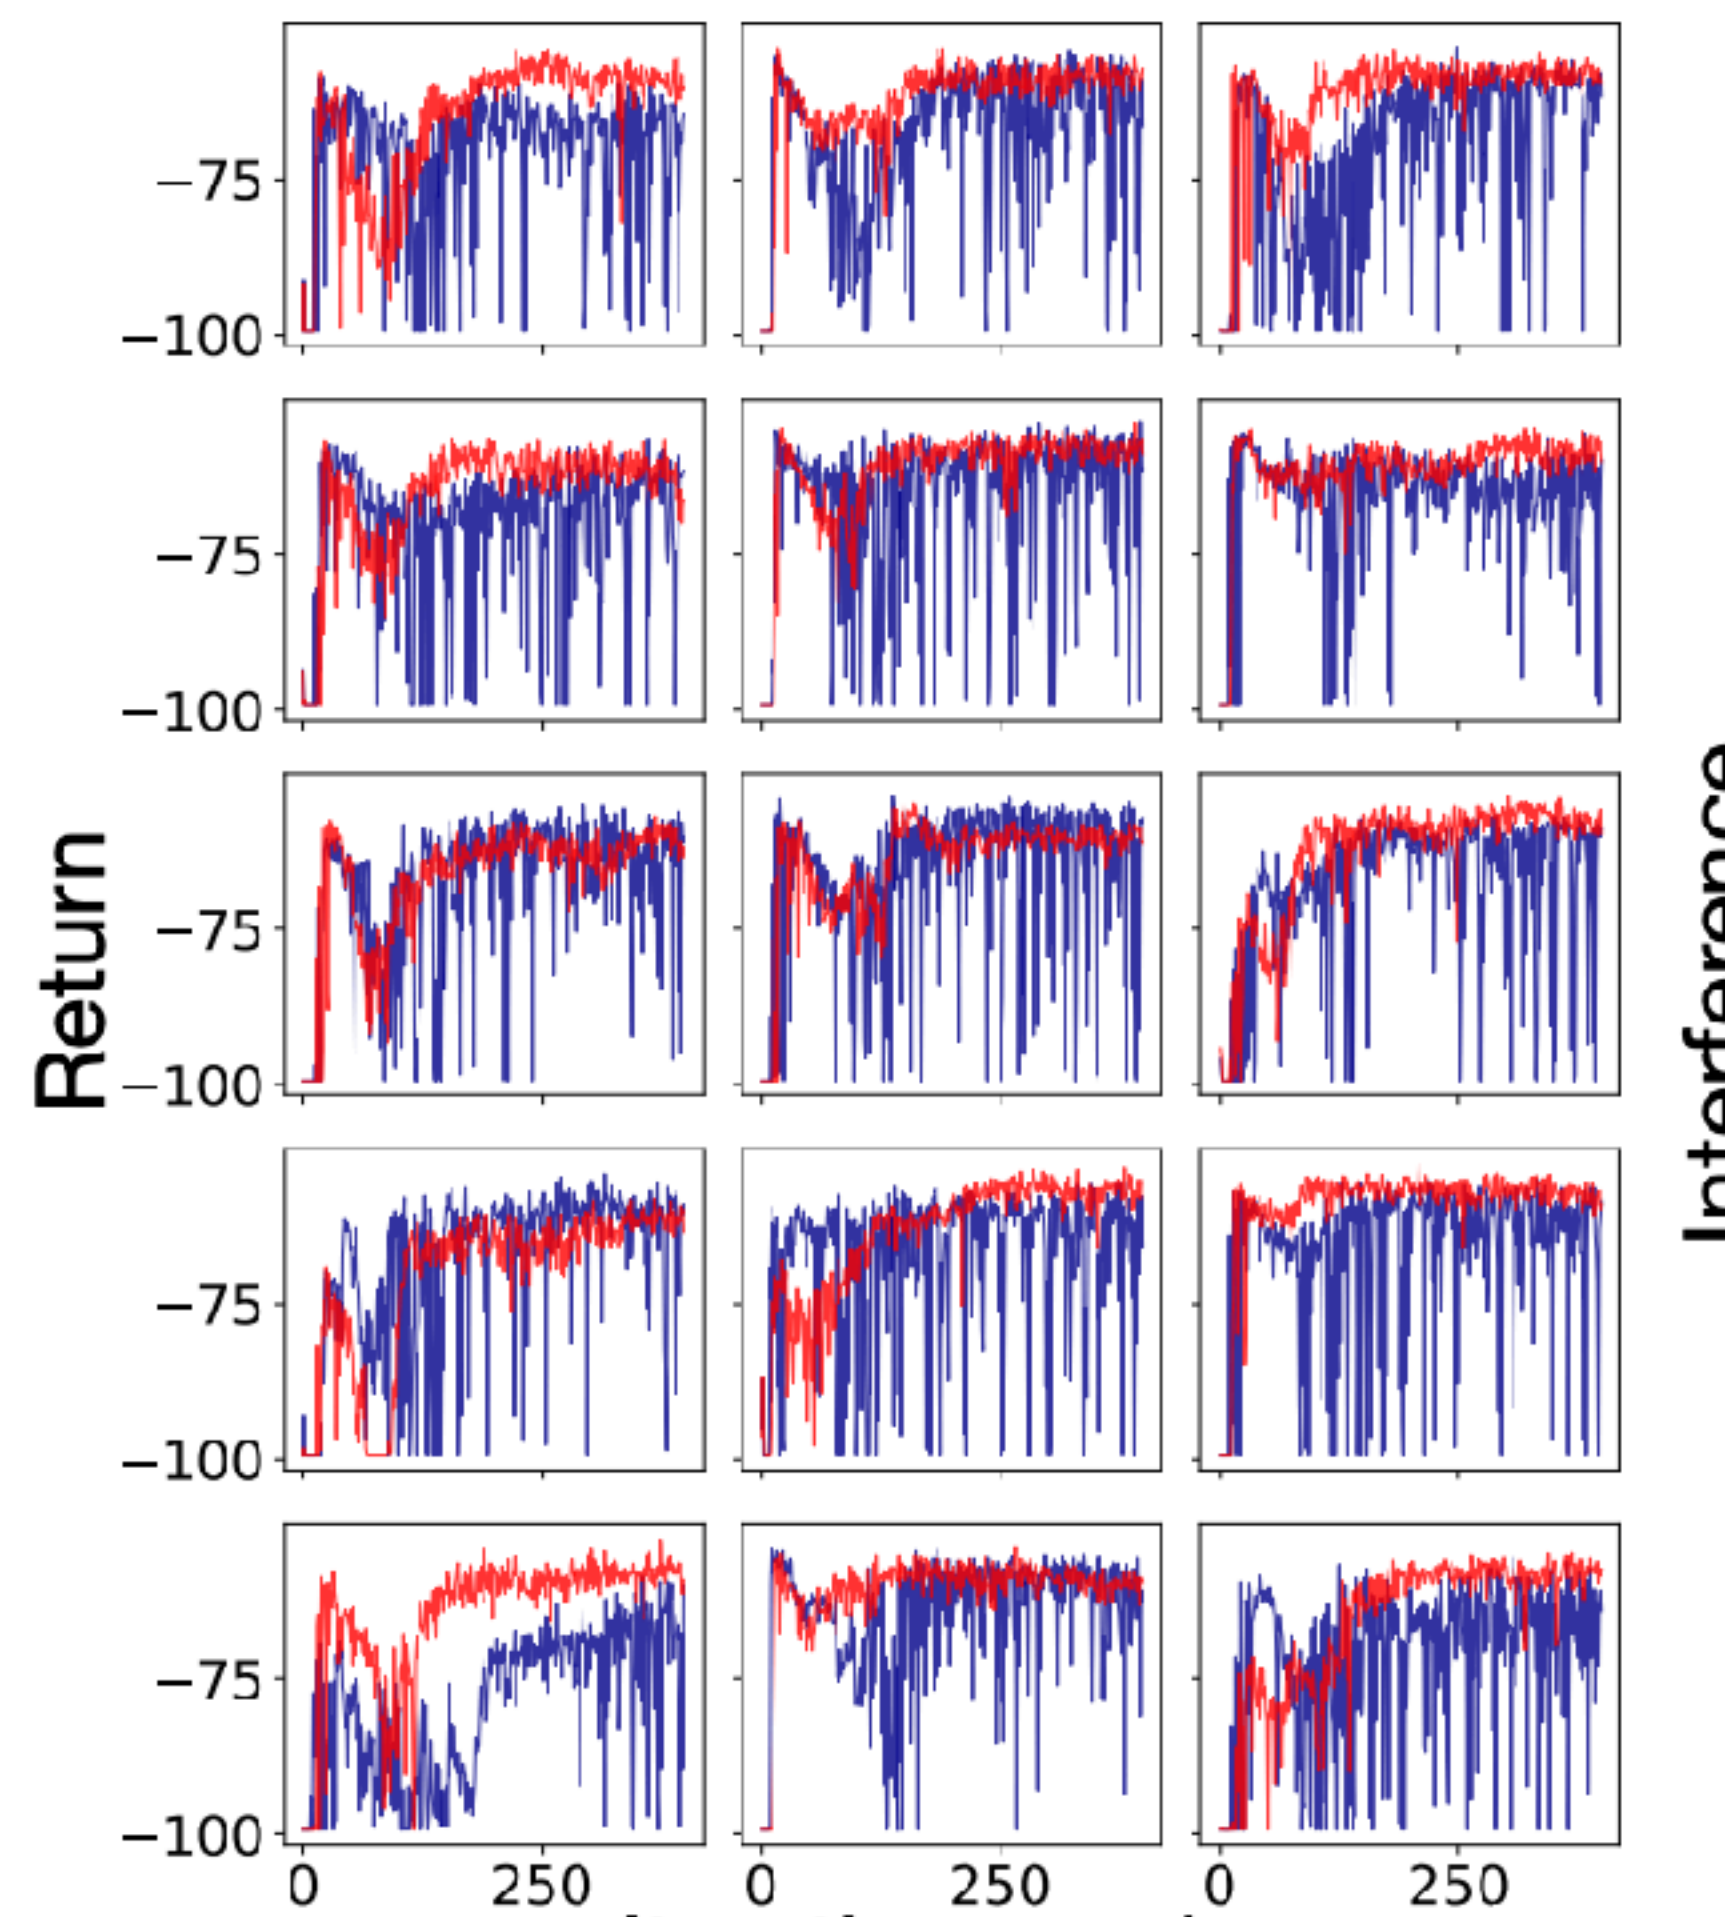
Example: sound off-policy control



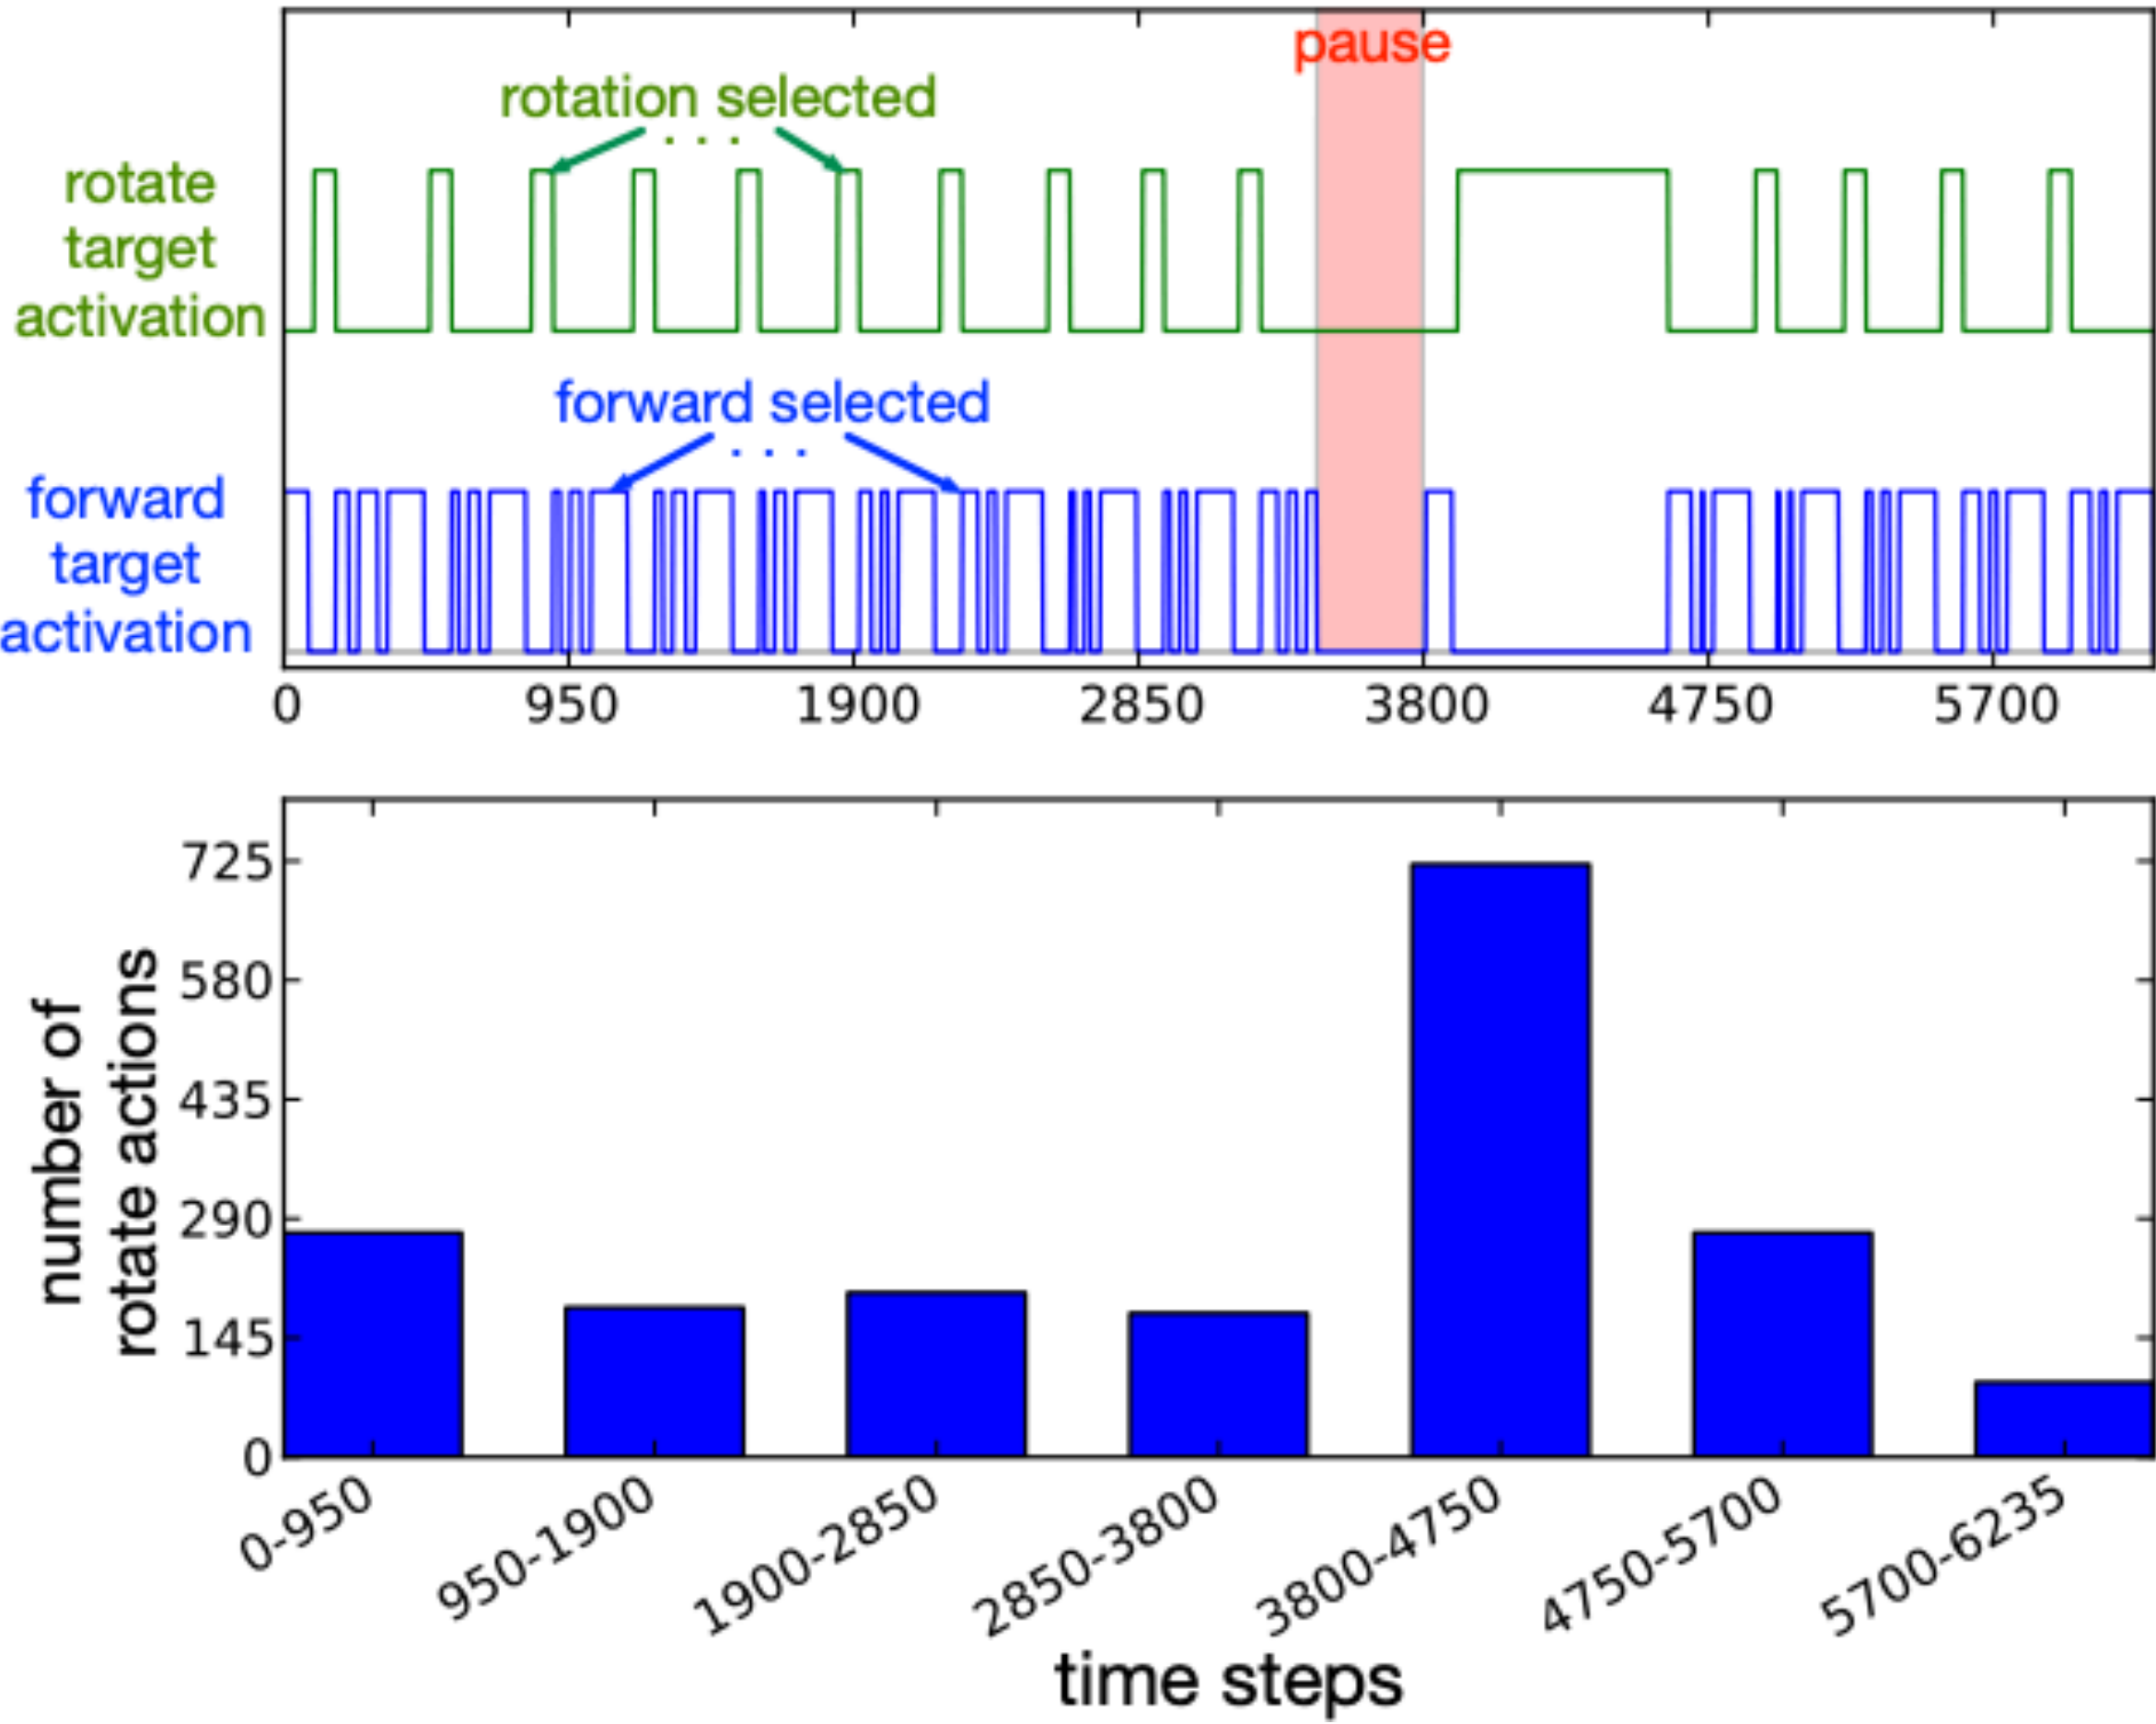
What to measure, what to plot?

- There are always multiple views into an experiment
 - There are many dimensions over which a new idea might be relevant
- This about what aspect is relevant to you and your problem:
 - Final value-function/policy quality/accuracy
 - Speed of learning
 - Insensitivity to hyperparameters
 - Robustness
 - Problem specific metrics
- Just in case: plot everything!

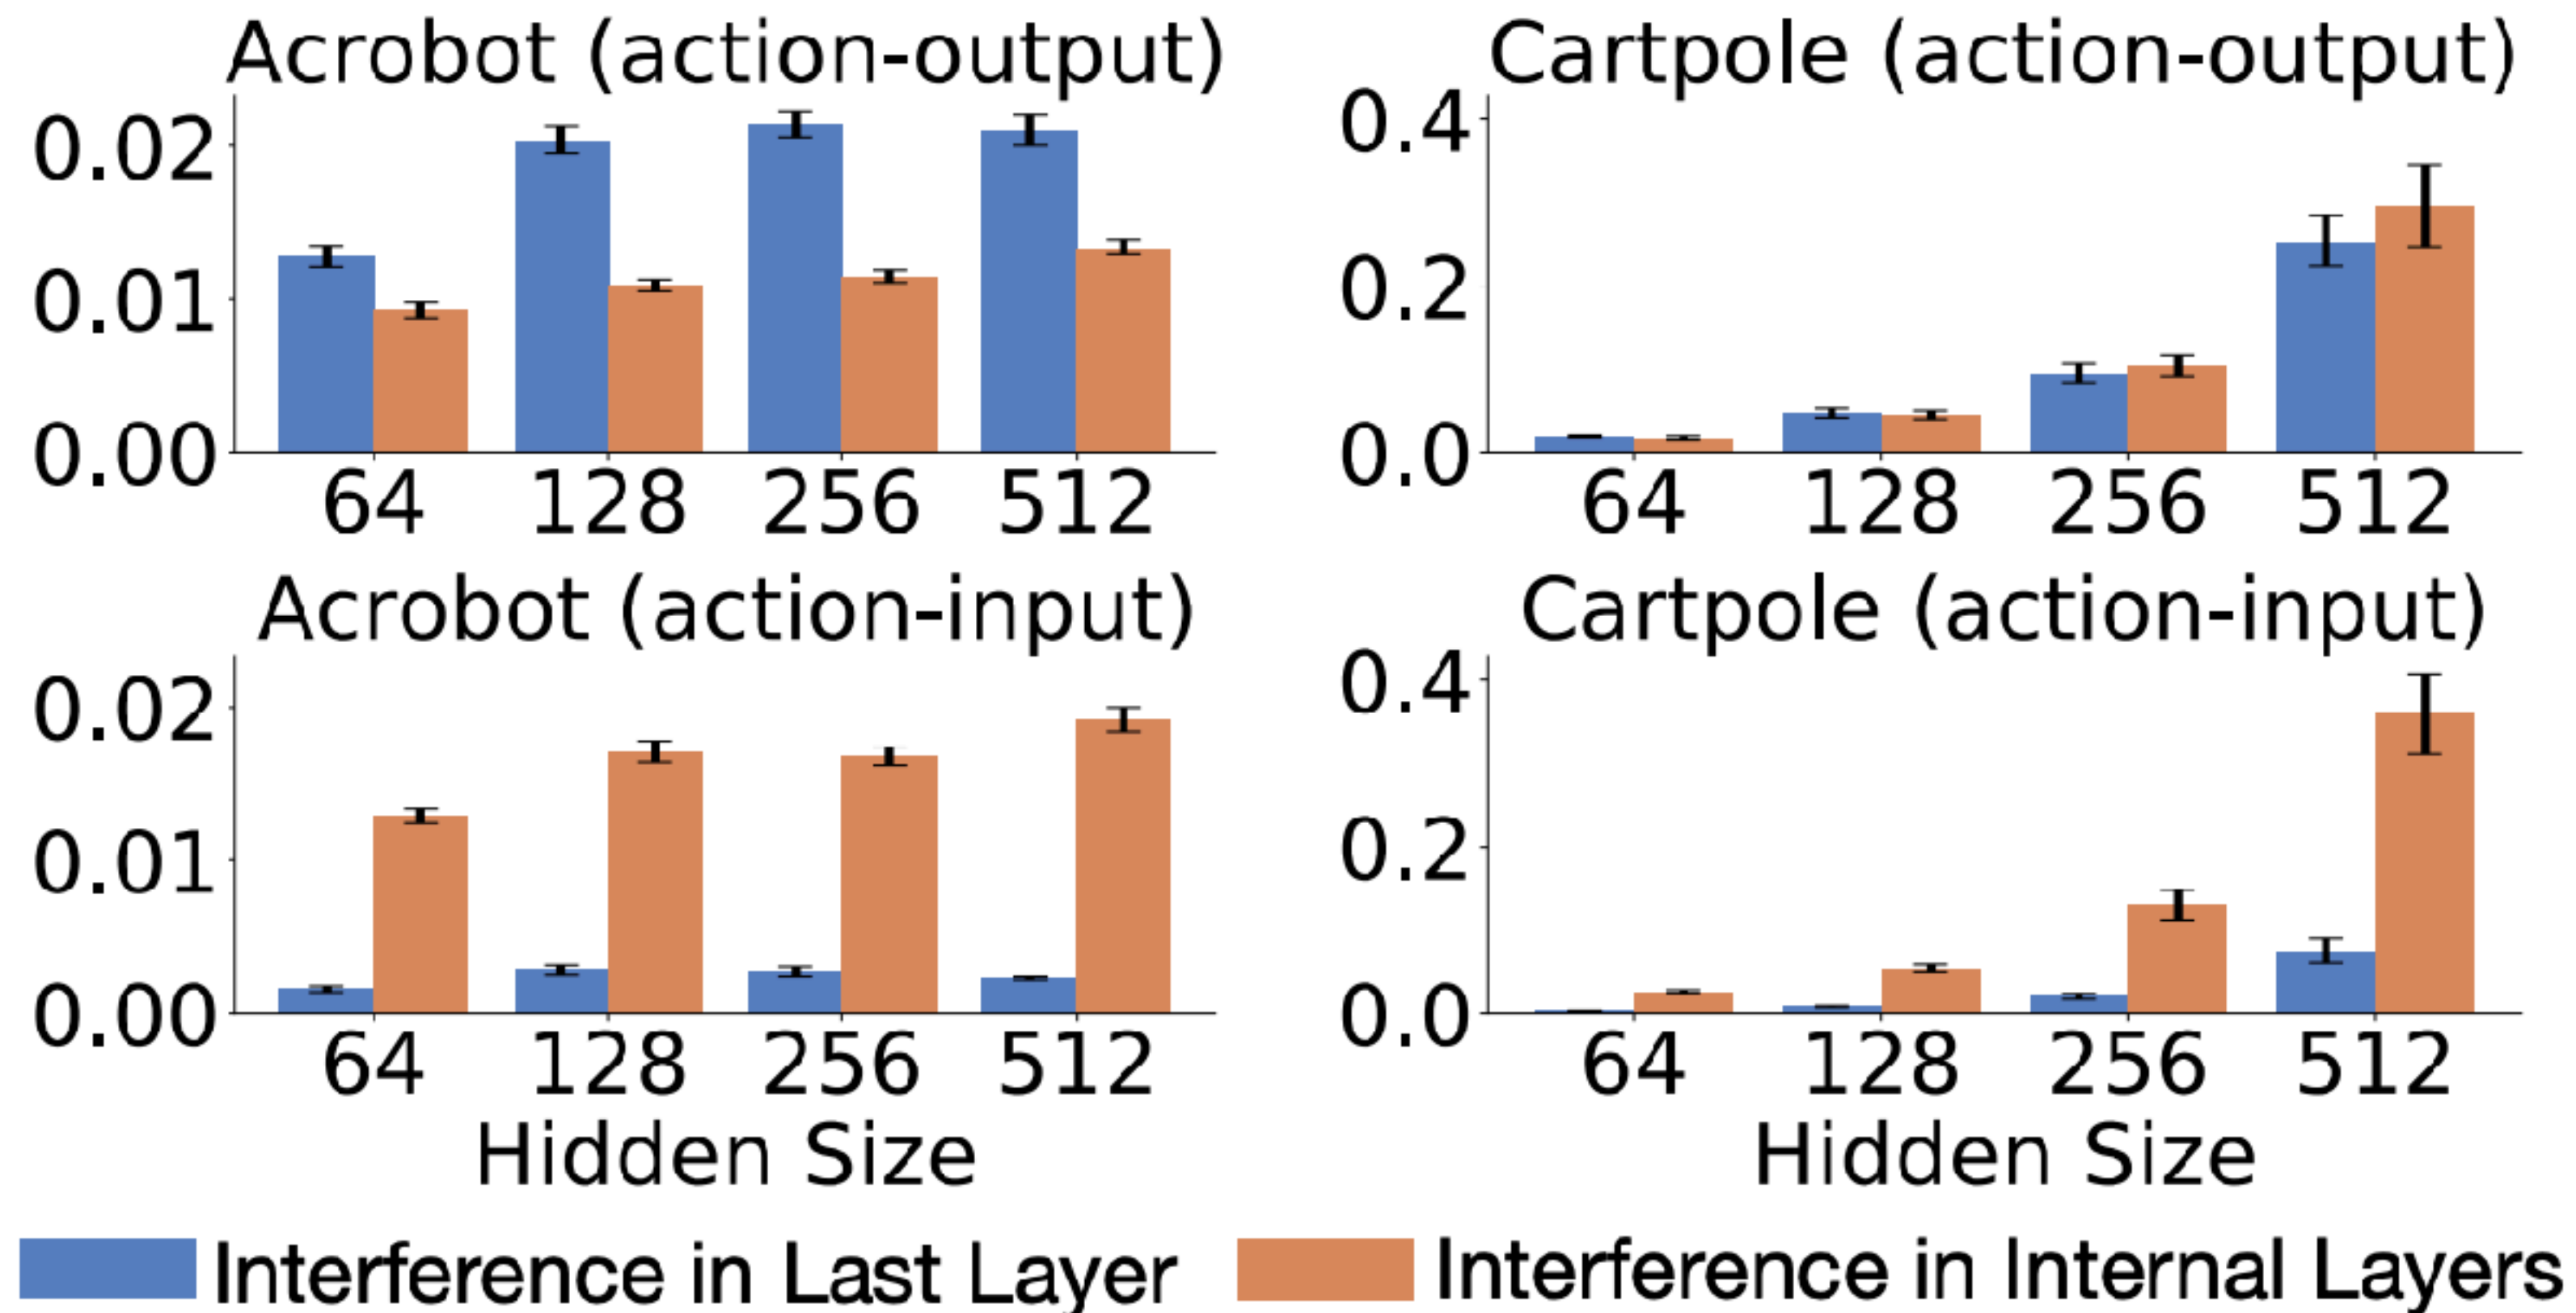
Example: a more stable control algorithm



Example: clear change in behavior



Example: where interference is happening in a network





Ultimately we end up comparing things

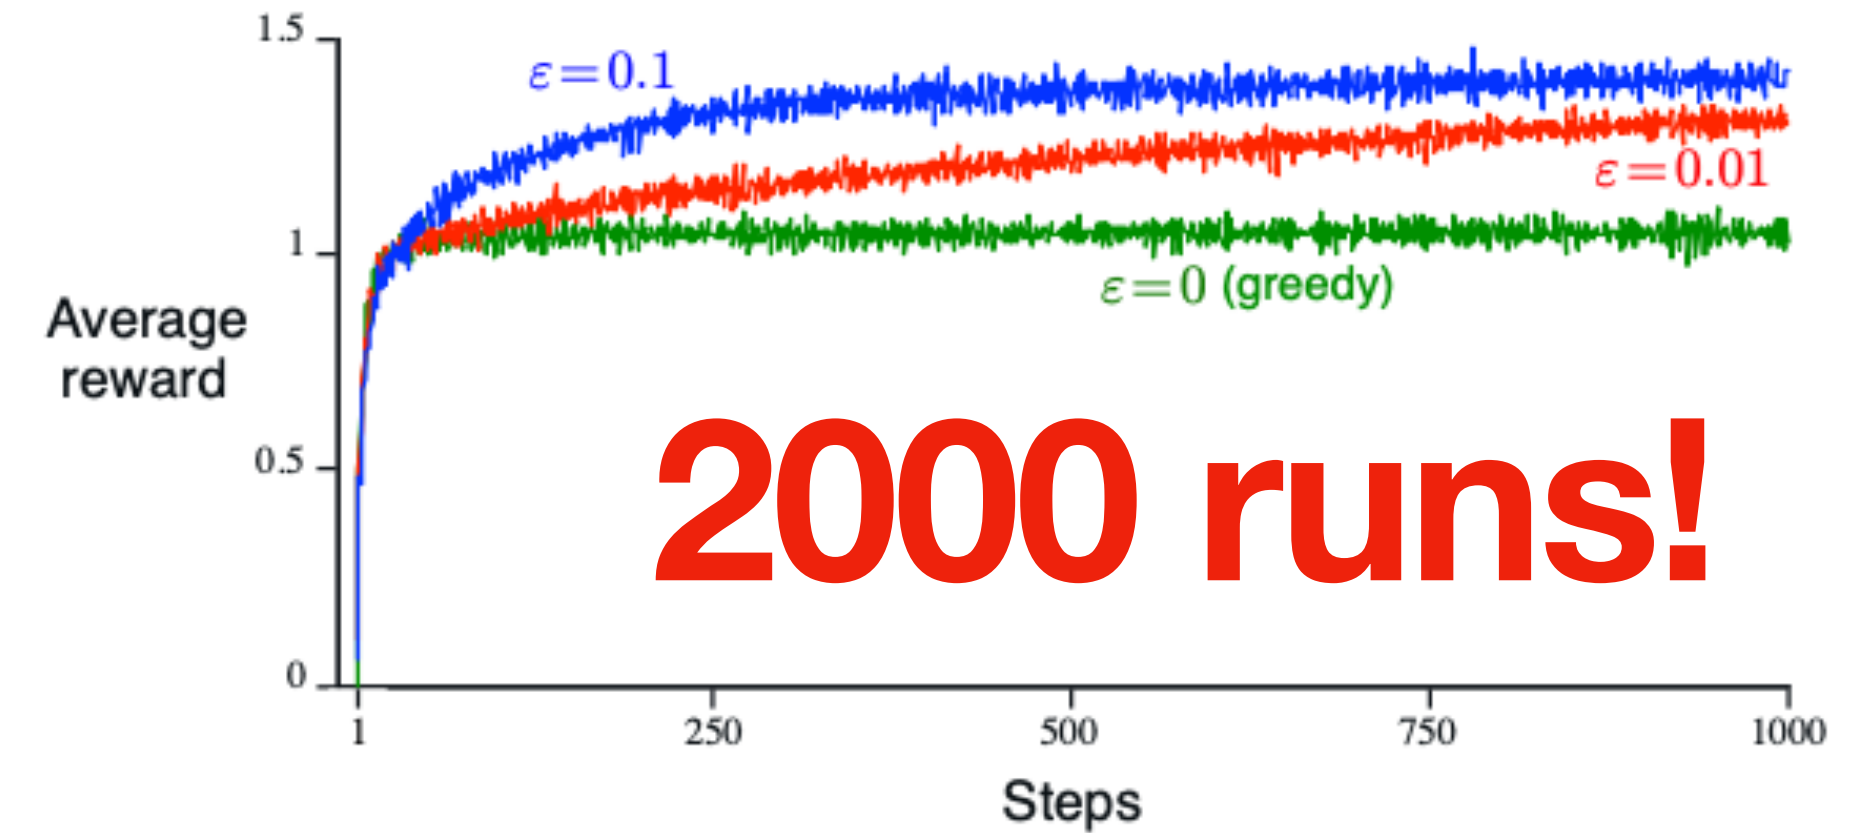
- SOTA competitor, natural baseline, or calibration agent
- We need to measure something & compare agents
- This is not about winning and losing ... its about telling the story of the data
- To tell the story accurately:
 - Properly report uncertainty & variation
 - Properly report how hard it was to get good performance
 - Properly report the impact of as many choices as you can
 - **Stretch:** properly reflect how well these algorithms might work in the real-world

Are our algorithms practically useful?

- **Mountain Car, Sarsa(λ) with tile coding — pretty much the best you can do on MC**
- Fixed start state, 0.5 decaying step size, 10 tilings 10x10





Without repetition we can say so little



- Experiment repetition is so important
- We don't want the results to be skewed by one algorithm getting lucky
 - Remember the MAB in Sutton&Barto...on some runs greedy is optimal
- We want to use statistical tools to talk about aggregate performance
- Hopefully we can build more reliable algorithms
- But we often need to look deeper to understand the mean & variance

The raw data can tell different stories




- 50 runs, 300 steps
- Credit: Andy Patterson



Which data/alg would you prefer?

Agents & Environments are data generators

- If we want to make statistical statements about the data, then we have to understand what it looks like
- **We want to turn a learning curve for a single run into a number**
- The first step is deciding on a measure of performance:
 - Total area under the learning curve (AUC)
 - AUC of the large x% of the data
- Other measures focused on stability are also possible but we will start with the classic ones



Getting one number



These are importantly different when sweeping hyper-parameters

The distribution of performance

- Given a set of AUC, one for each run, what does the distribution of those numbers look like?
 - Bell shaped / Normal /Gaussian
 - Skewed
 - Multi-modal
 - Flat or point mass?
- **Practical tip:** set the seed for the environment and the agent independently, and use the run number for reproducibility
- What should we do about the hyper parameters?

The distribution of performance


- Given a set of AUC, one for each run, what does the distribution of those numbers look like?
 - Bell shaped / Normal /Gaussian
 - Skewed
 - Multi-modal
 - Flat or point mass?
- **Practical tip:** set the seed for the environment and the agent independently, and use the run number for reproducibility
- What should we do about the hyper parameters?

How many runs do we need?

- Common practice is 3, 5, ... maybe 8
- In the literature you can find up to thousands of runs
- Let's run an experiment:
 - Mountain Car with random starts
 - Sarsa(λ) with tile coding — reasonable hyper parameter choices
 - We will plot mean episodic return over 250 episodes
- What story does the data tell?


What if we did 3 runs?

- Histogram of mean episodic return over 100k steps (around 250 episodes)




What if we did 5 runs?

- Histogram of mean episodic return over 100k steps (around 250 episodes)




Many runs are needed to see the shape of the distribution

- Histogram of mean episodic return over 100k steps (around 250 episodes)



- Estimating the agent's performance accurately requires many independent repetitions of the experiment


Environments design choices matter too




- Notice how the distribution was a bit skewed, not perfectly bell shaped
- We can get other distribution shapes by including **cutoffs**:
 - Restarting the episode if the agent reaches a max number of steps
 - This ensures the no episodes a really bad—**might make bad agents look good**
 - This gives free exploration—especially if random starting states are used

Cut-offs skew performance

Regular MC->




MC w cut-offs->




- 1000 step episode max

Every agent & environment pair can be different

- Same experiment and setup in Puddle world:




- Mountain car:




Design choices interact

- Mountain car with two different start states:

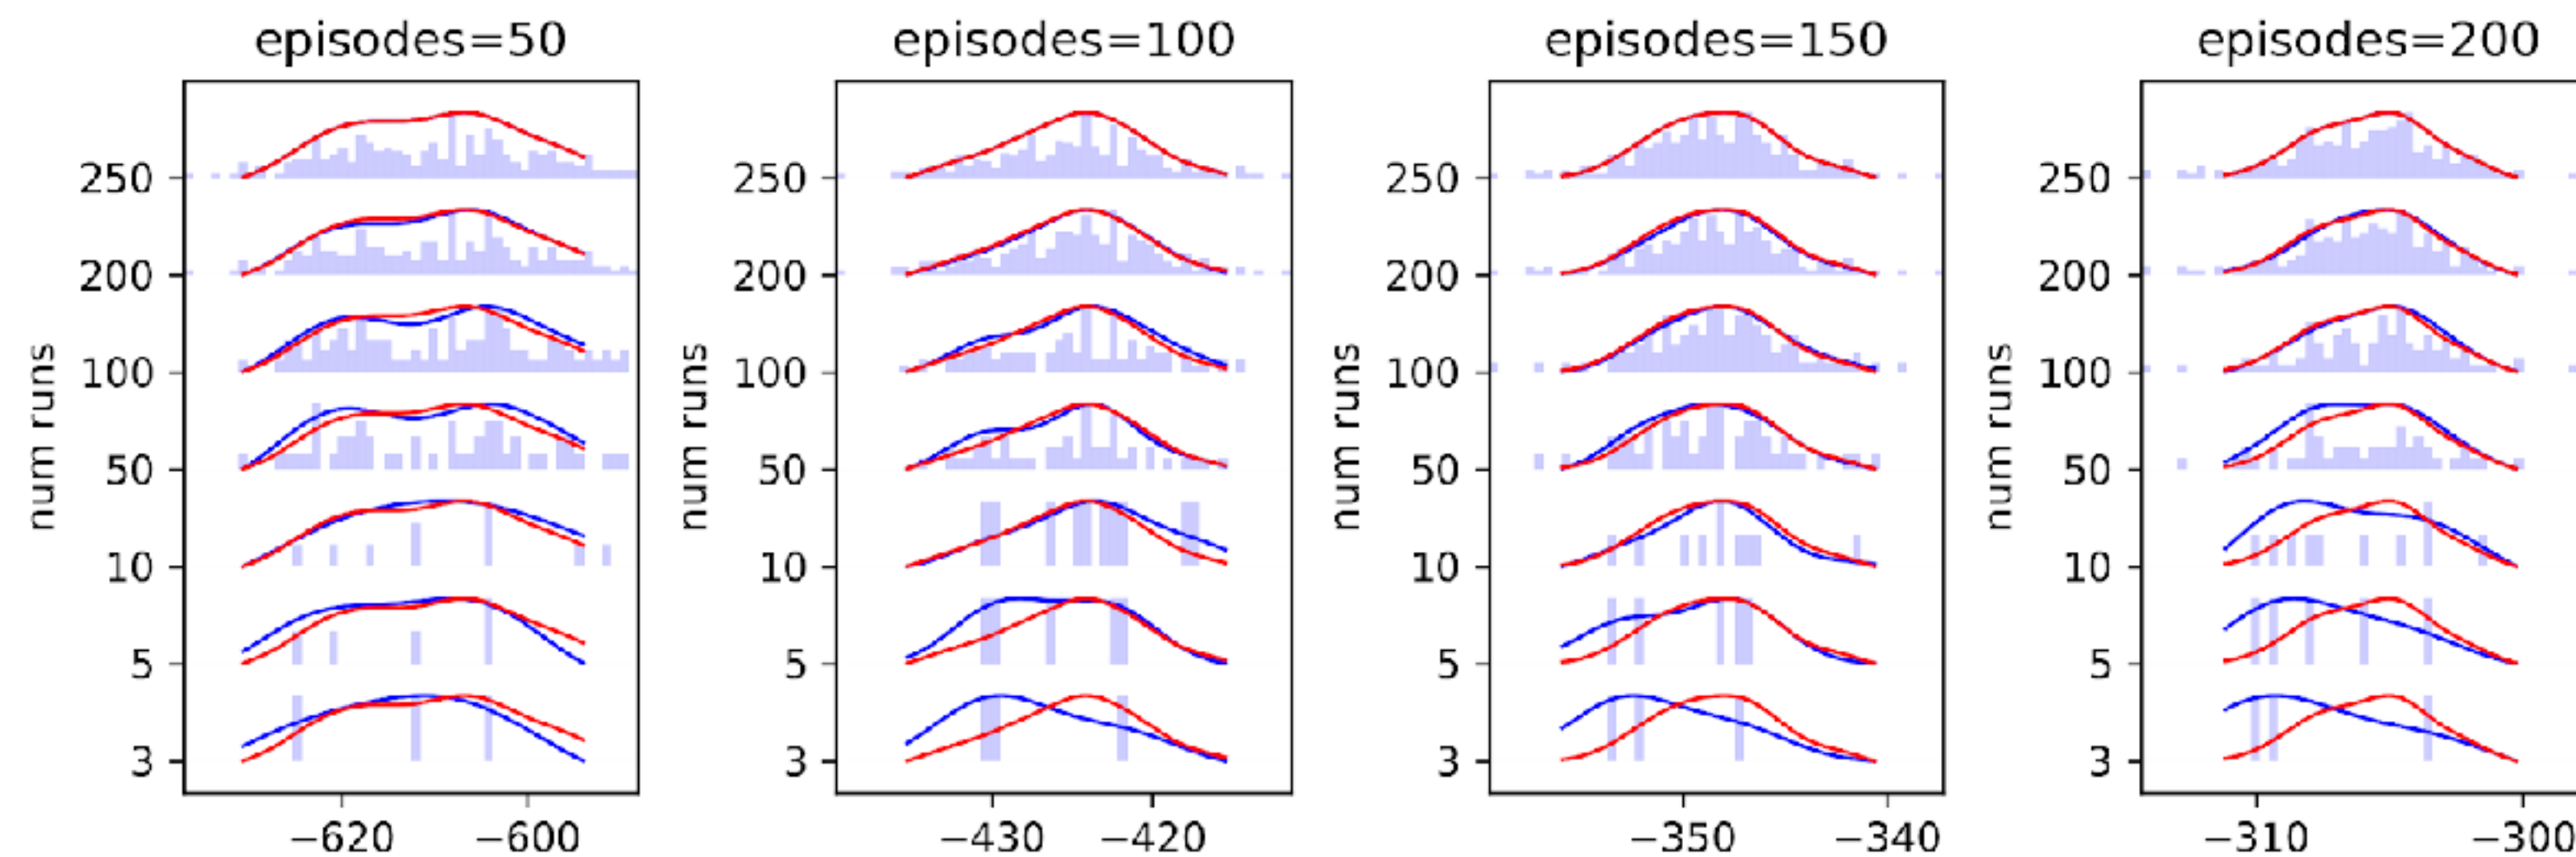


- Mountain car with two start states and cutoffs:

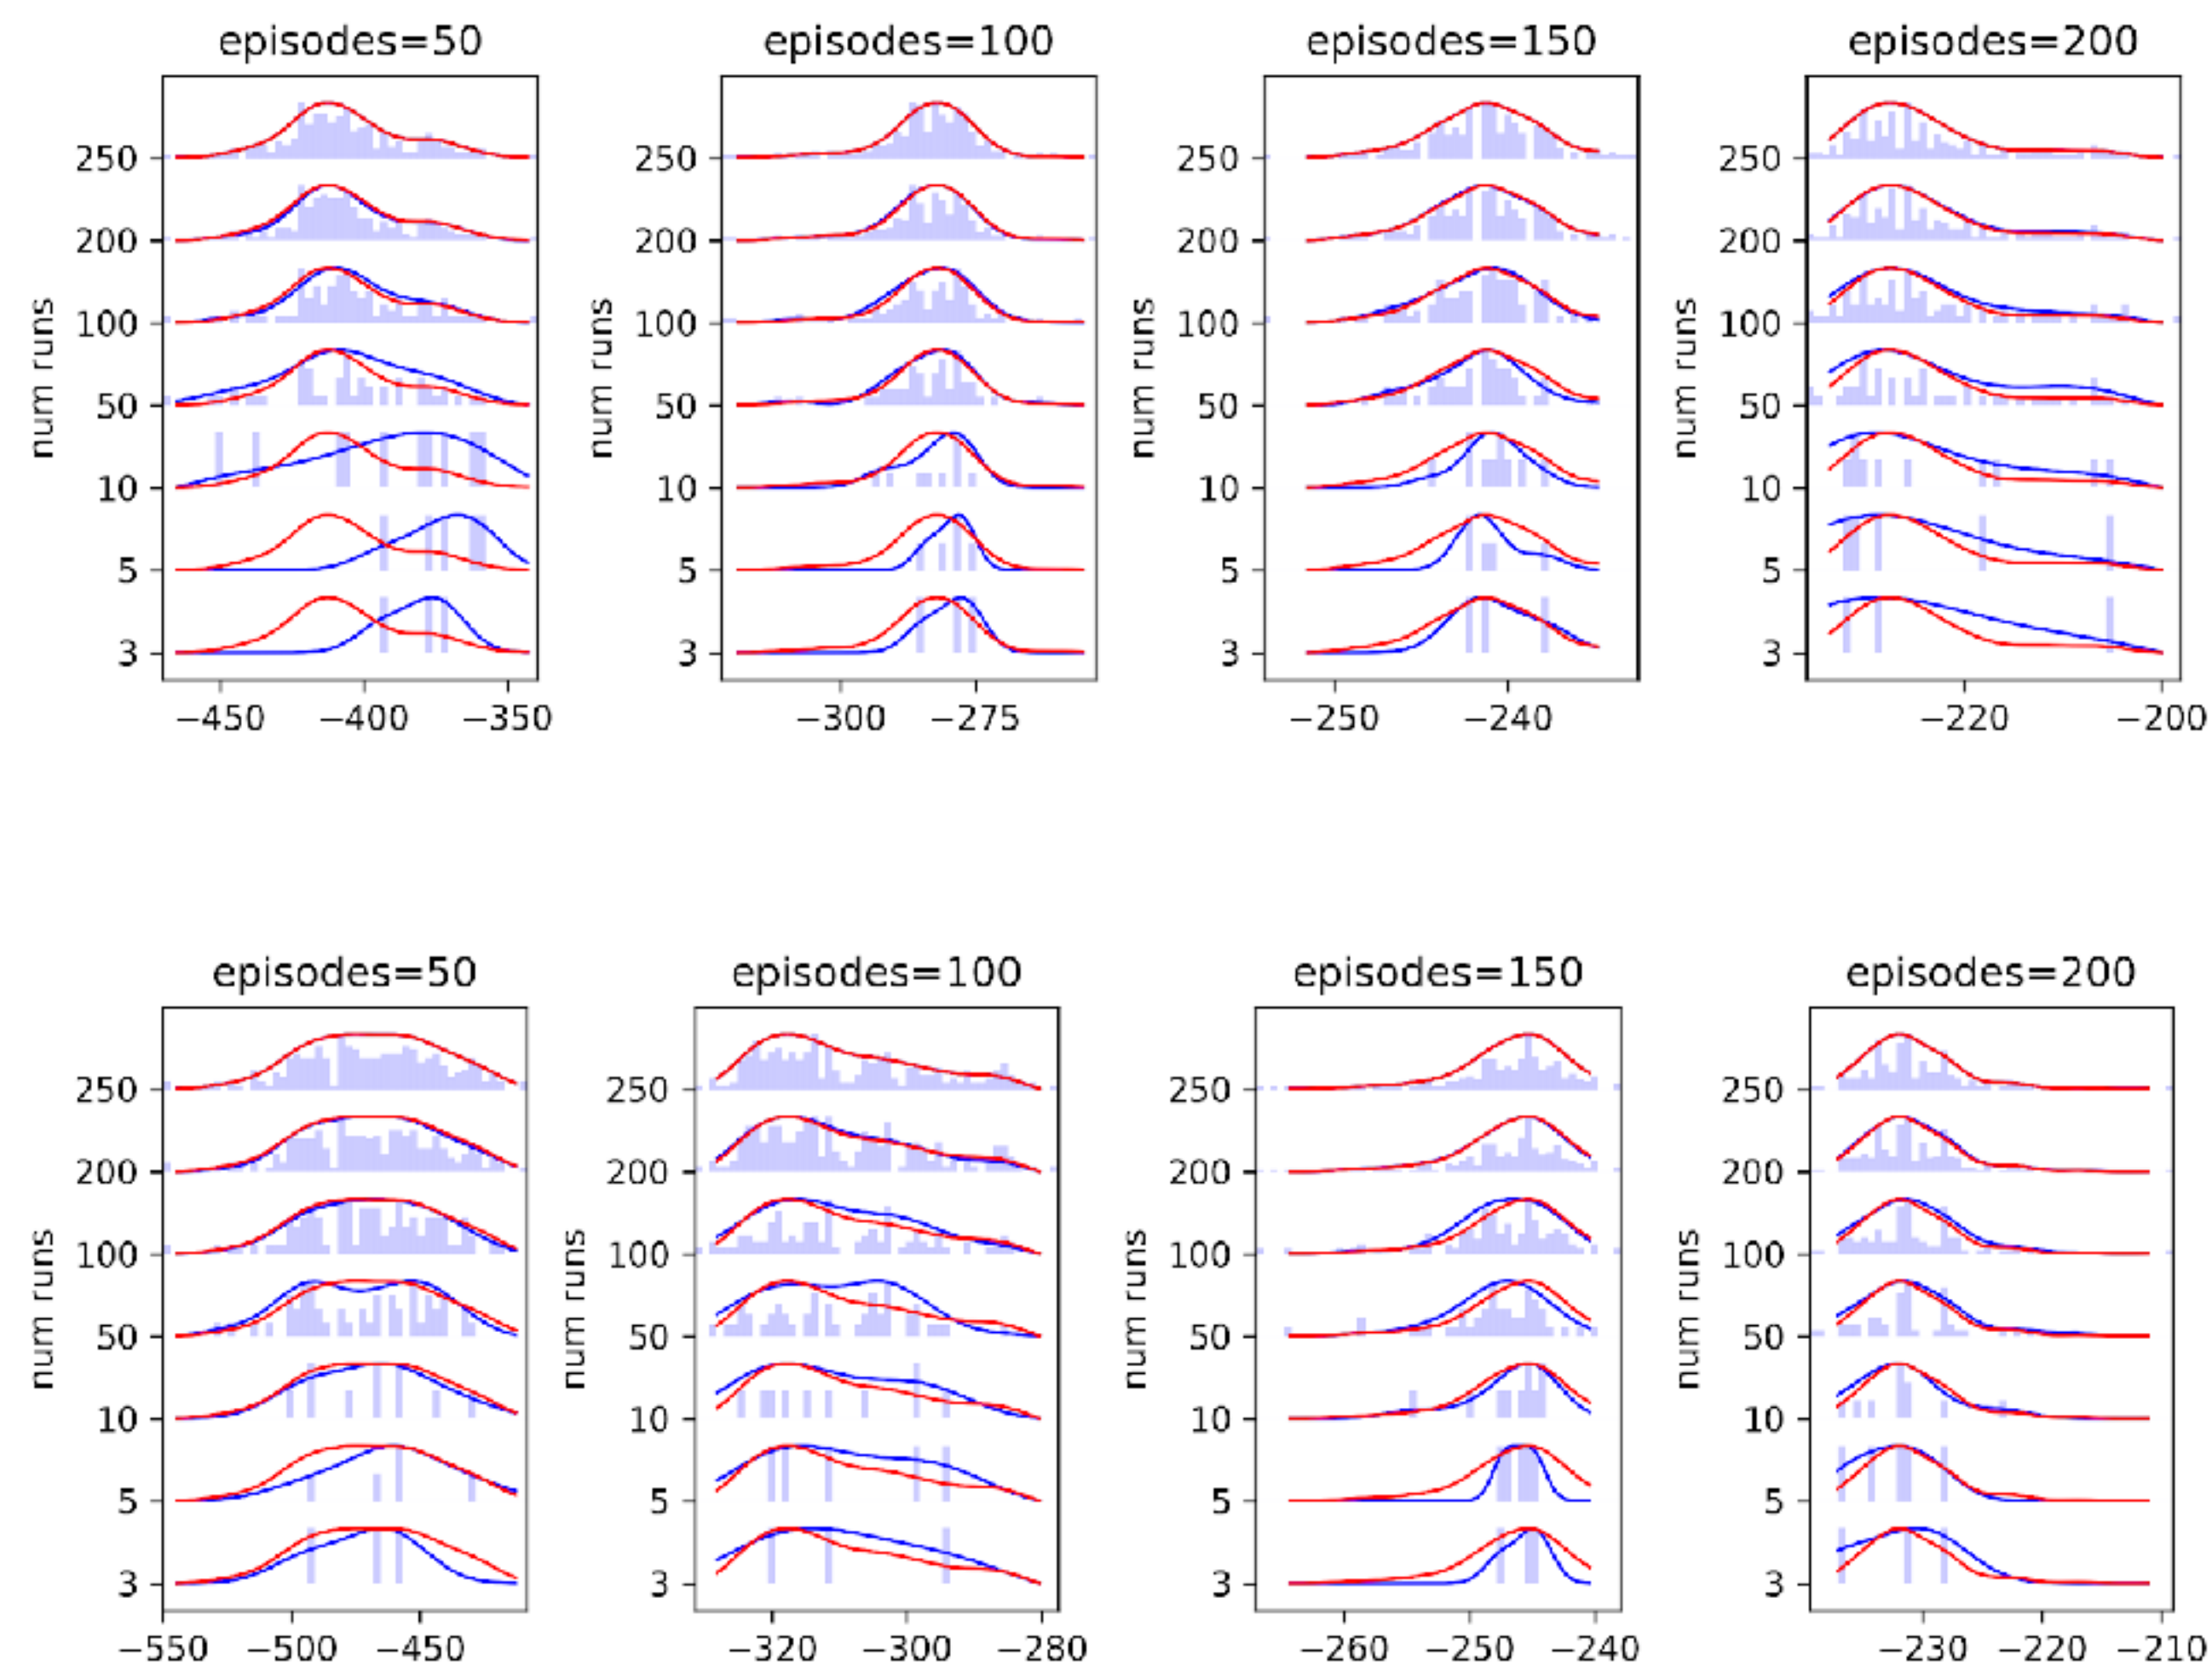


Experiment design choices interact too

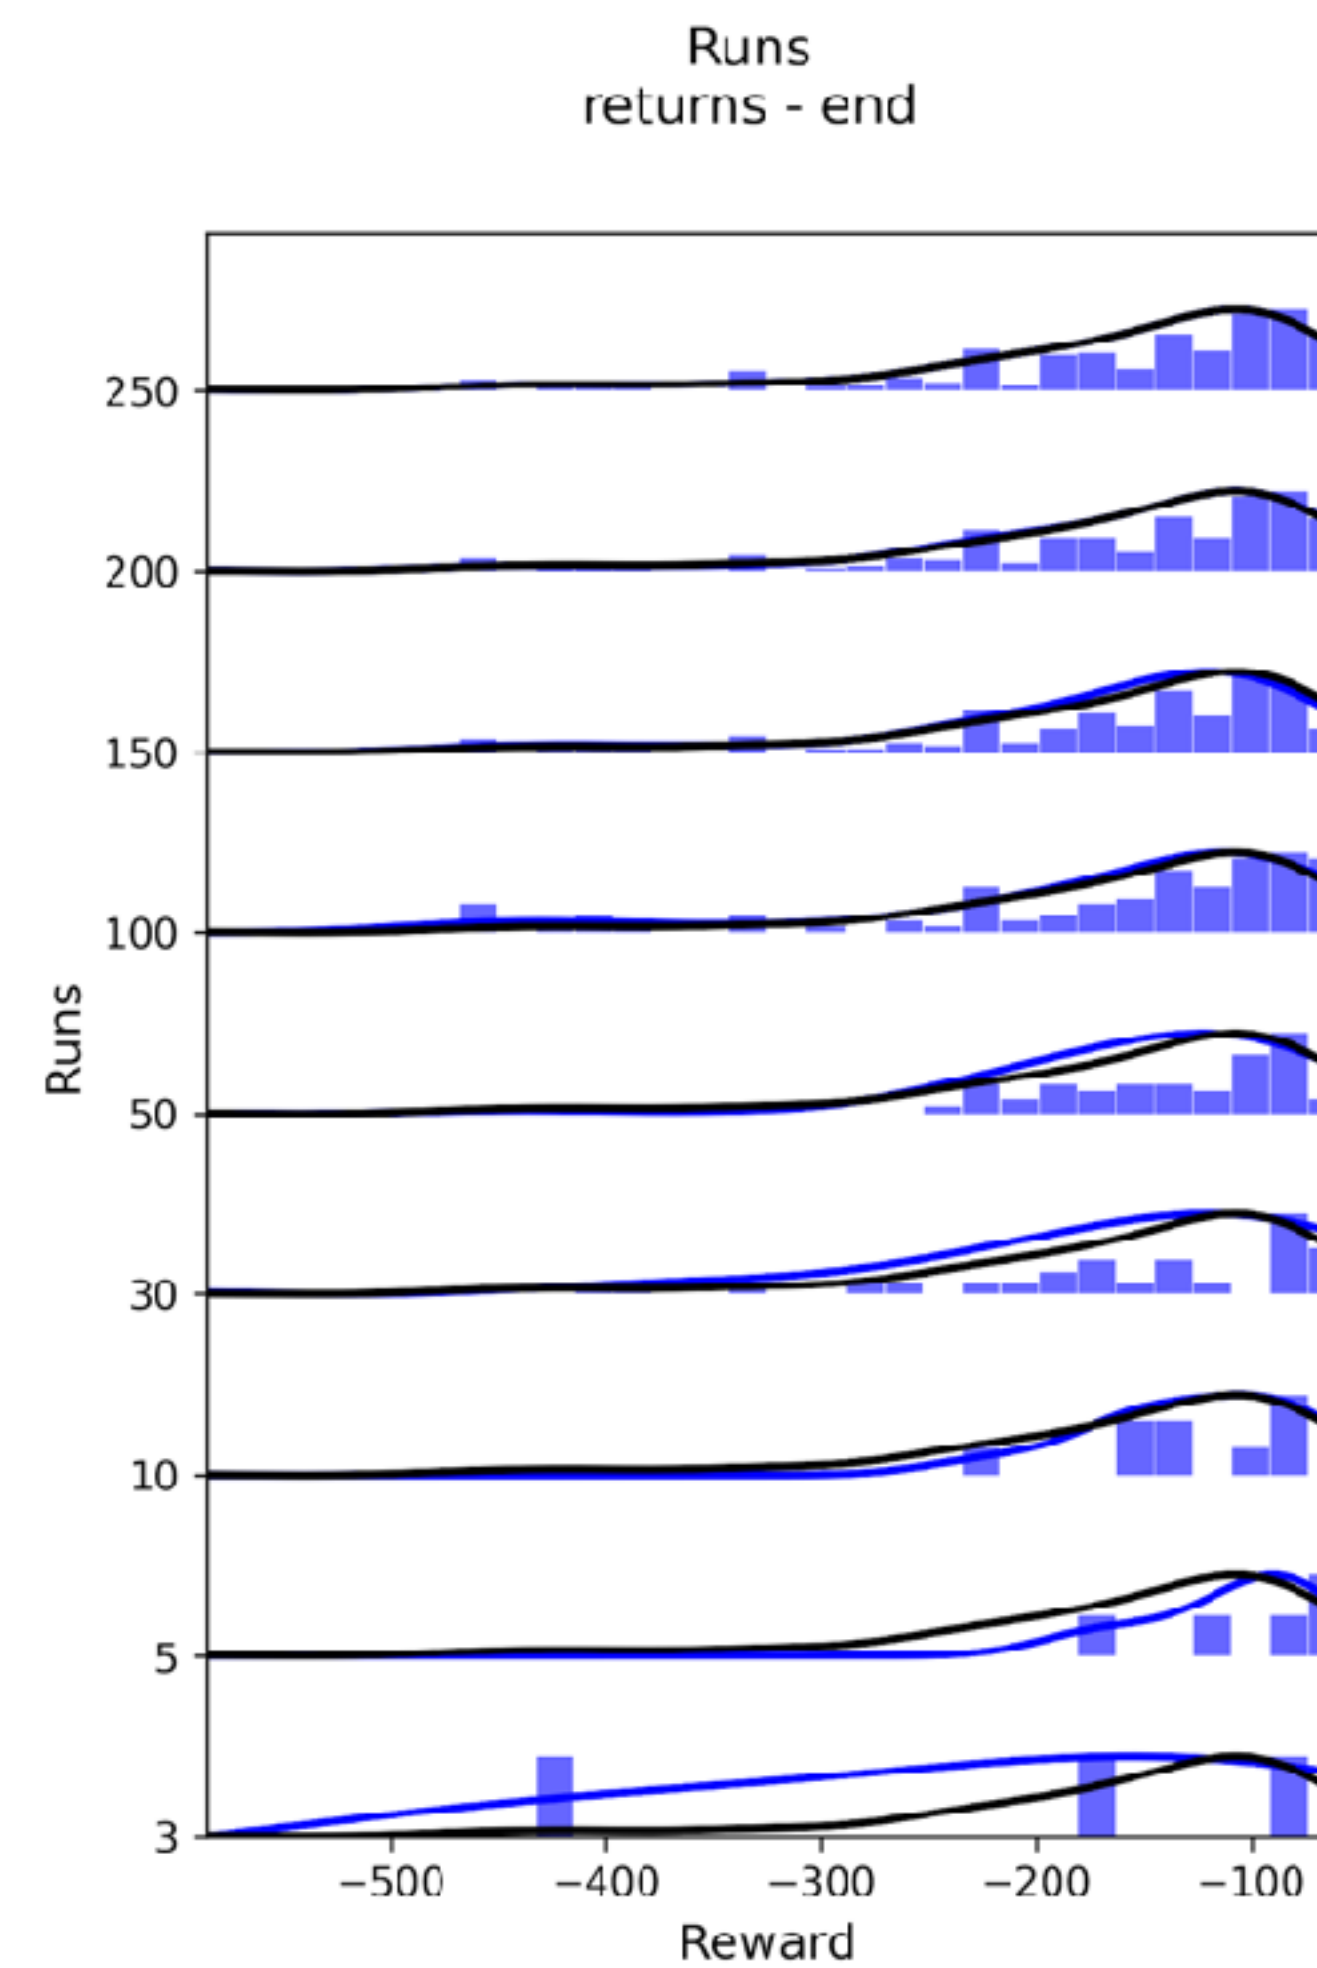

- In the prior plots we always ran 100k steps, and looked at the dist with more and more runs
- We can also look at the dist with more and more episodes (MC) ...





With and without cut-offs



In puddle world we see impact of performance metric




A closer look at cut-offs in puddle world



Bi-modality can even happen without explicit effort

MountainCar - mellowmax

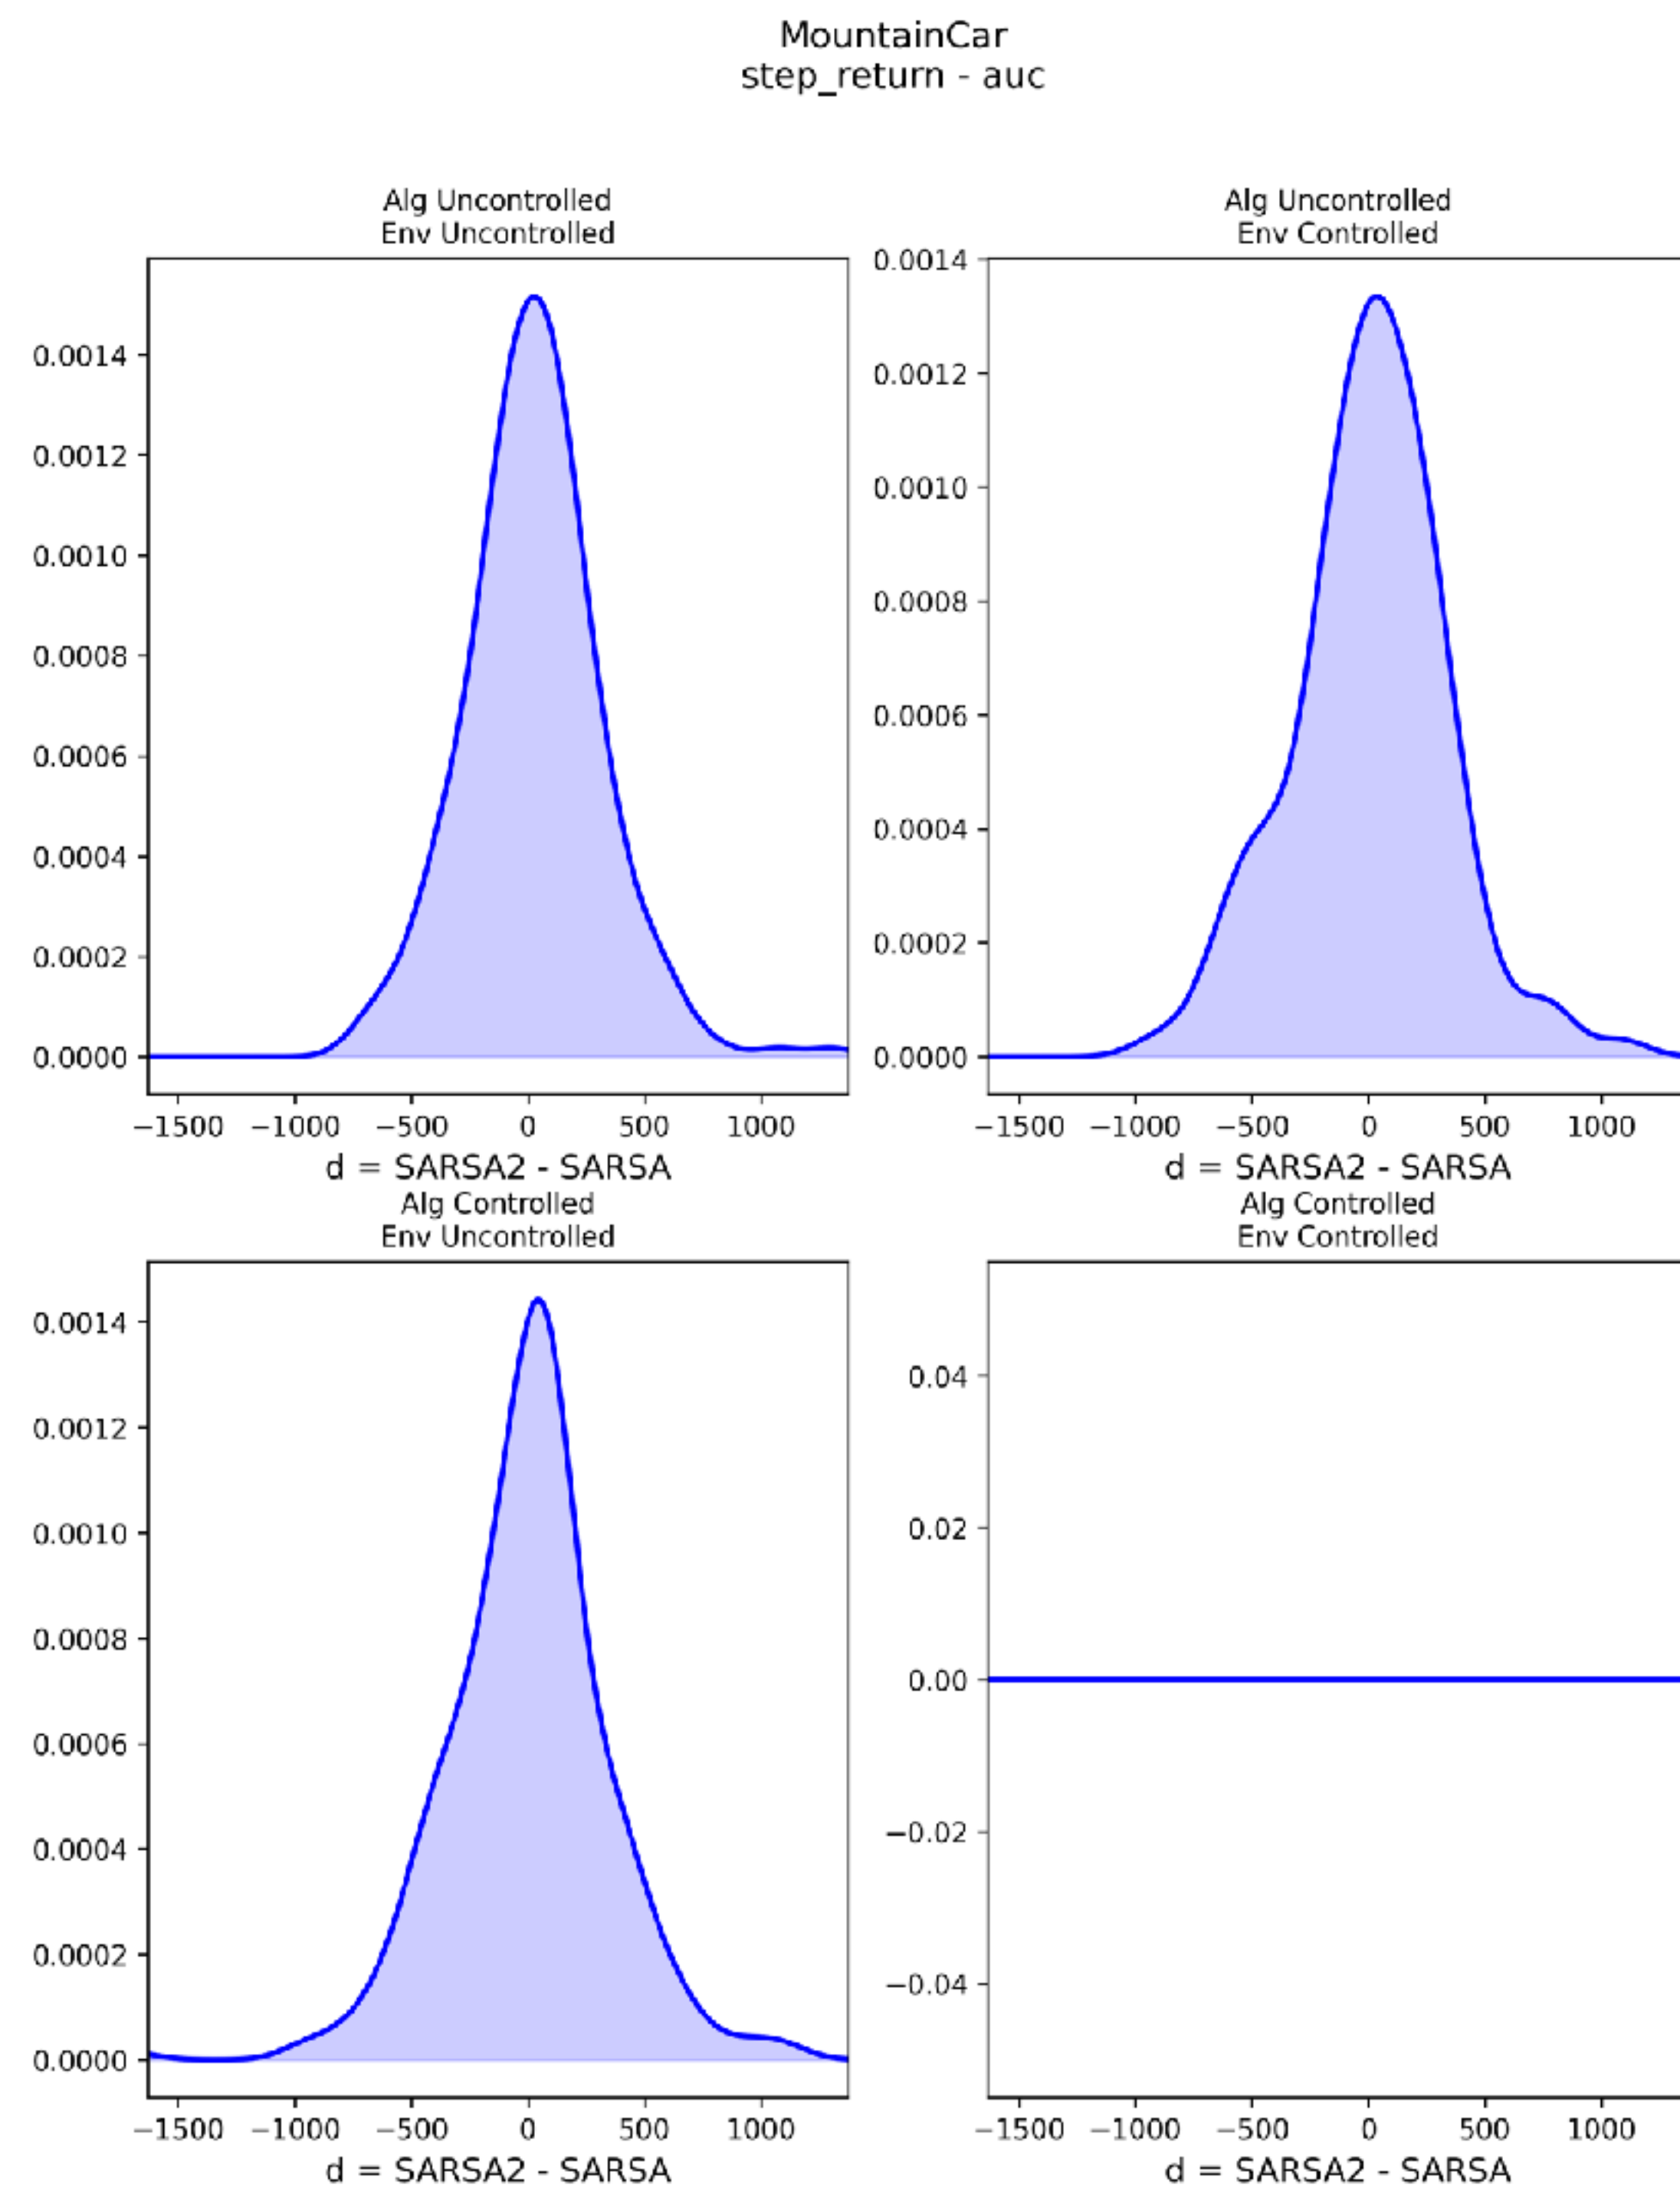


- Mountain car with 3 different algorithms and a Neural Network (2 layer, 32 hidden units, relu)
- Max episode length=1000, 100k steps total
- Agent hypers:
 - epsilon=0.1
 - Adam with beta_1 = 0.9 and beta_2=0.999
 - buffer_size = 4000, batch_size=32
 - No target nets

Controlling randomness

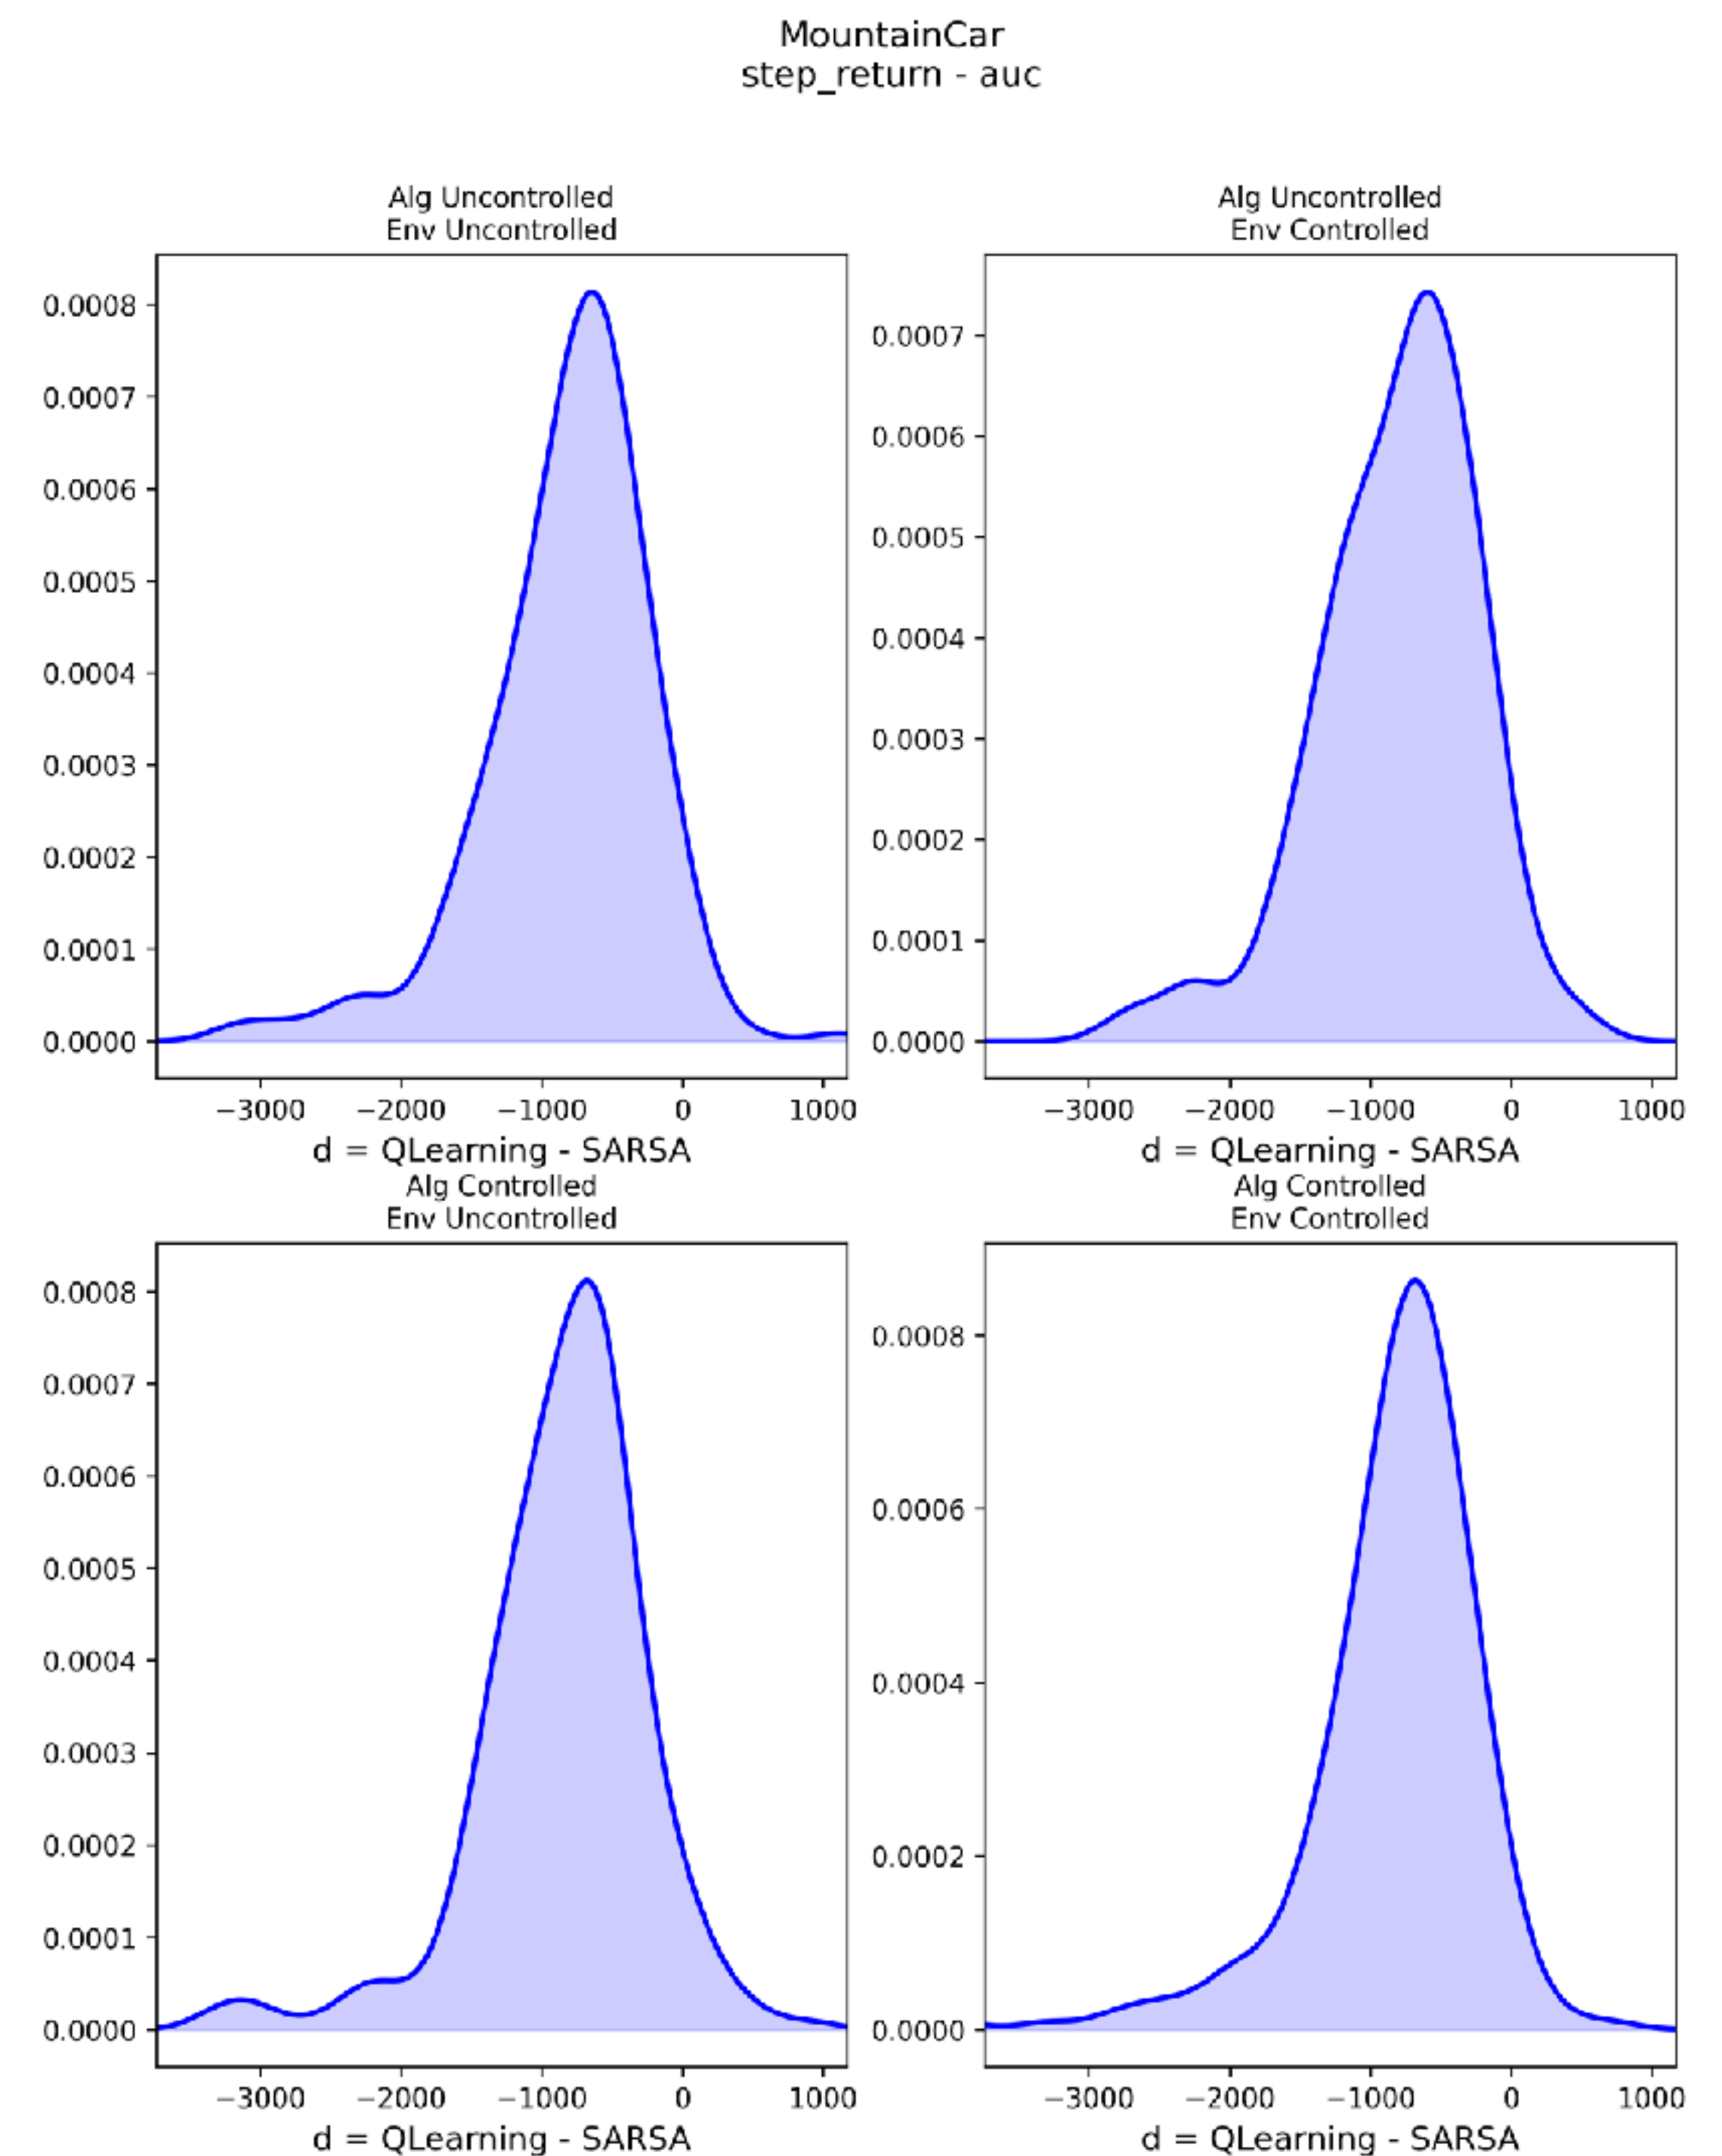
- Typically both the agent and environment have different sources of randomness:
 - In mountain car the start states, and epsilon in the agent for example
- We can decide to control these sources of randomness or not:
 - Controlled means the seed to the agent/env random number generator is set with the run_number
- There are 4 possibilities for controlling and not controlling each

Controlling randomness: comparing the same algorithm (250 runs)



Controlling randomness: comparing Q-learning and Sarsa


Sarsa > Qlearning here



Controlling randomness: comparing Q-learning and Sarsa

but with only 5 runs

Qlearning looks better than before —>



Why it all matters

- We can't always show all the data
- Worse: depending on experiment, environment, and agent design choices the data will all be different
- We will be left with mountains of data; dozens of plots
- That's no fun for us, and certainly no good for a paper
- We want to aggregate the data, and use statistical tools like hypothesis tests and confidence intervals to make broader conclusions

**You can't just compute
error bars and report
p-values blindly**

Hypothesis testing

- Let's say we draw samples from two populations, with true means μ_0 and μ_1
- We estimate the mean of each population: \bar{x}_0 , \bar{x}_1
- Then we want to determine if the populations have different means
- We use a hypothesis test:
 - Null hypothesis: $\mu_0 - \mu_1 = 0$ (the true means are the same)
 - Alternative hypothesis: $\mu_0 - \mu_1 \neq 0$ (the true means differ)
 - **We want to reject the null hypothesis!**

How probable is it to observe this sample or a more extreme one, given that there is no true difference in the performances of both algorithms?

The p-value is that probability: to reject the null we want it to be extremely unlikely that we observe differences in the sample means given that the algorithms indeed perform the same!


If your p-value is large, then your evidence (data) does not provide enough support to reject the null

Hypothesis testing

- Let X_1 be the random variable denoting the performance of algorithm_1
- Let X_2 be the random variable denoting the performance of algorithm_2
- If we assume X_1 and X_2 are normally distributed
 - Therefore $X_{\text{diff}} = X_1 - X_2$ is normally distributed
- We want many samples of X_{diff} (say 30 or more)

Hypothesis testing procedure

- **Let** $X_{diff,1}, X_{diff,2}, \dots$ be a sequence of RV representing runs of the experiment and \bar{X}_{diff} = average of $X_{diff,1:n}$
- **True distribution** over the differences: $\bar{X}_{diff} \sim p_{true}$, i.e., $p(\bar{x}_{diff})$ is density
- **Sample** $\bar{x}_{diff,0}$ // we run an experiment
- **Assume** null hypothesis: p_{null} is defined such that $\mathbb{E}[\bar{X}_{diff}] = 0$
 - This is the hypothesized model of p_{true}
 - E.g., p_{null} might be a mean-zero Gaussian over \bar{x}_{diff}
- **Question:** how likely is $\bar{x}_{diff,0}$ under H_0
i.e., how likely is it that we would see $\bar{x}_{diff,0}$ or a more extreme value:
 $p_{null}(\bar{X}_{diff} > \bar{x}_{diff})$ (if unlikely, then our model likely wrong)



Is the difference significant?

A difference is called significant at significance level $\alpha/2$ when the p-value is lower than $\alpha/2$

Key assumptions in hypothesis testing

- We most often use a t-test (and standard error bars)
- ~~They assume the distributions of performance are Normal~~
- Performance is measured at random and independently from one another (each agent)
- Same sample size
- Continuous and bounded performance distributions
- ~~Equal standard deviations~~